

Librarian

Project memory for coding agents

A record of real use

Keep the architecture, decisions and fixes learned in one session available to the next. Librarian turns those findings into a wiki that is checked against your source code.

Librarian is a memory layer for coding agents. It saves useful findings as project notes, retrieves them through MCP or a CLI, and checks them against source code on a schedule. Set it up once, then let each session build on the last.

2,638

Memory calls

500

Findings recorded

54

Active days

How I have used it

I have used this memory workflow with Claude Code and Codex since 17 June (I added logging to analyse how it is doing later). The record covers nearly three months, with activity in July, August, September and October.

The wider workflow contains 39,904 tool calls, including 4,436 identifiable MCP calls. The memory server recorded the 2,638 calls shown above. These are separate views of usage, so their counts are not added together.

The habit that makes memory useful

Recall: Find the relevant note before exploring the repository.

Verify: Check the source files needed for the current change.

Improve: Record what changed. The scheduled Librarian audits it.


Sources: vault transcript export and server log. Static transcript extraction can miss calls inside batches.

Three months of project knowledge

Knowledge that builds with each investigation

By the end of July, the log contained 187 findings. By 8 October, it contained 500. Each useful addition gives a later session something it can reuse.

Recorded findings



Cumulative remember events. These are proposals, not a count of unique facts. July and October are partial months.

Period	All tools	Memory	Total findings	Log to date
Jul 13-31	10,094	875	187	299 KiB
August	13,092	1,077	315	628 KiB
September	10,800	360	400	765 KiB
Oct 1-8	5,918	326	500	900 KiB

The memory calls include 1,302 note reads and 674 searches. Agents also recorded 500 findings, listed notes 132 times and checked pending updates 30 times.

The vault snapshot contains 325 project notes. The event history includes 105 knowledge corrections and 122 gotchas. Log sizes are cumulative reserialized UTF-8 JSONL. Monthly call counts vary with workload; they are not productivity scores.

The economics of reuse

How repeat work can save 30 percent

A known answer can replace another round of discovery. The saving comes from reaching the right source files faster while keeping the checks needed for the change.

~207

Tokens per median note

970


Requests revisit a note

Measured: the median of 1,160 note bodies is 826 characters (~207 tokens at 4 chars/token). Of 1,302 requests, 970 revisit a note after project aliases are normalized.

Task input budget	No memory	With Librarian
Locate the known answer	3,250	0
Task work and source checks	4,750	4,750
Read the note and maintain it	0	207 + 400
Build cost over 20 reuses	0	4,400 / 20
Total token equivalents	8,000	5,577

2,423 fewer tokens

$(8,000 - 5,577) / 8,000 = 30.3\%$



What must be true. This example assumes an 8,000-token task and that the note replaces all 3,250 locating tokens. The prior local audit estimated those locating tokens and a 4,400-token build cost per note. Ongoing upkeep is budgeted at 400 tokens per reuse.

For 30%, total memory overhead must stay below 850 tokens. This model uses 827. Use a known note path: the prior audit estimated a broad search at 4,235 tokens. These are conditional input-volume savings, not measured billed savings. Source: memory-layer-efficiency-forensics.md and MCP server log.

Librarian and a Similar repo (Agent Memory)

Two ways to keep project knowledge

Both keep linked memory in Git. Agent Memory Repo offers a flexible format and a skill. Librarian packages a workflow for coding agents, with structured capture and scheduled source audits.

Workflow	Agent Memory Repo	Librarian
Recall	Read MEMORY.md, grep notes or follow links.	Search project notes through MCP or the CLI. Read a specific note when its path is known.
Capture	Agents edit notes directly and commit changes.	Agents append separate proposal files. The Librarian merges accepted findings into the wiki.
Maintain	The spec describes periodic Dreaming to merge, clean up and check sources. The trial adds no scheduler.	A packaged scheduler runs source audits. Note size and link checks keep the wiki navigable.
Set up	Invoke the skill. The trial is local; sharing uses persistent storage or a private remote.	Init wires supported agent clients and schedules the Librarian. Add registers each repository.
A useful fit	Personal and team memory, including saved queries, scripts and linked context.	Repeated coding work across repositories where decisions and fixes need ongoing source verification.

Why I use Librarian

I want findings from today's investigation to be available next time, with a routine that checks them as the code changes. The benefit depends on agents using that knowledge to narrow their work.

Try it on a repository you revisit

github.com/Dbhardwaj99/librarian

Sources: project READMEs and skills, checked 8 Oct 2026. No head-to-head token benchmark.